

✓ MIS373 - AI For Business - Assignment 2

Task 1: House Price Prediction

Student Name: VABI SABHARWAL

Student ID: S223036676

Table of Content

1. [Executive Summary](#)
2. [Data Preprocessing](#)
3. [AI Model Development](#)
4. [Experiments Report](#)

✓ 1. Executive Summary

Use this section to introduce the business problem, data set, method, experiments, and obtained results

Business Problem:

The property market experiences frequent changes, creating difficulties for buyers and sellers to accurately assess property values. In response to this challenge, our objective was to create a predictive model capable of estimating house prices based on pertinent features. Such a model has the potential to support stakeholders in making well-informed decisions regarding property transactions.

Dataset:

We used a detailed dataset with info on residential properties like location, size, rooms, amenities, and past sale prices. This dataset was crucial for training and testing predictive models, offering a wide range of features to understand property values better.

Method:

Our strategy involved utilizing machine learning methods, specifically regression algorithms, to construct predictive models capable of estimating house prices. We conducted experiments with various models, including linear regression, k-nearest neighbors (KNN), random forest, and grid search with cross-validation (GridSearchCV), to determine the most optimal approach for our objective.

Experiments:

During our experiments, we divided the dataset into two parts: training and validation sets. The models were trained using the training set and then assessed based on their performance on the validation set. To gauge the accuracy of the models, we utilized the root mean squared error (RMSE), a widely-used metric in regression tasks that indicates the difference between predicted and actual values.

✓ Obtained Results:

Our experiments yielded the following results:

1. Model 1 (Linear Regression): Validation RMSE of 645,596
 2. Model 2 (KNN): Validation RMSE of 206,717
 3. Model 3 (Random Forest): Validation RMSE of 207,433
 4. Model 4 (GridSearchCV): Validation RMSE of 216,778
 5. Model 5 (Random Forest with hyperparameter tuning): Validation RMSE of 159,318
-

✓ 2. Data Preprocessing

Carry out necessary data preprocessing and exploration.

IMPORTING THE LIBRARIES

```
from __future__ import print_function
import os
import math
import datetime
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
```

```
from sklearn.model_selection import cross_val_score
from sklearn.model_selection import KFold
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, MinMaxScaler
from sklearn.preprocessing import OneHotEncoder
from sklearn.impute import SimpleImputer
from sklearn.pipeline import Pipeline
from sklearn.metrics import mean_absolute_error
```

Adjusting Display Options

```
pd.set_option('display.max_columns', None)
pd.set_option('display.max_rows', None)
```

IMPORTING THE DATASET

```
house_data = pd.read_csv("Part1_house_price.csv")
house_data.set_index('id', inplace=True)
house_data.head(10)
print('Number of records read: ', house_data.size)
```

Number of records read: 400000

Displaying the First Few Rows of the House Data

```
print(house_data.head())
```

	id	date	price	bedrooms	bathrooms	sqft_living	\
	7129300520	20141013T000000	221900.0	3	1.00	1180	
	6414100192	20141209T000000	538000.0	3	2.25	2570	
	5631500400	20150225T000000	180000.0	2	1.00	770	
	2487200875	20141209T000000	604000.0	4	3.00	1960	
	1954400510	20150218T000000	510000.0	3	2.00	1680	

	id	sqft_lot	floors	waterfront	view	condition	grade	sqft_abov
	7129300520	5650	1.0	0	0	3	7	118
	6414100192	7242	2.0	0	0	3	7	217
	5631500400	10000	1.0	0	0	3	6	77
	2487200875	5000	1.0	0	0	5	7	105
	1954400510	8080	1.0	0	0	3	8	168

	id	sqft_basement	yr_built	yr_renovated	zipcode	lat	lon
	7129300520	0	1955	0	98178	47.5112	-122.25
	6414100192	400	1951	1991	98125	47.7210	-122.31
	5631500400	0	1933	0	98028	47.7379	-122.23
	2487200875	910	1965	0	98136	47.5208	-122.39
	1954400510	0	1987	0	98074	47.6168	-122.04

	id	sqft_living15	sqft_lot15
	7129300520	1340	5650
	6414100192	1690	7639
	5631500400	2720	8062
	2487200875	1360	5000
	1954400510	1800	7503

DATA TYPES

house_data.dtypes

```
date           object
price          float64
bedrooms       int64
bathrooms      float64
sqft_living    int64
sqft_lot       int64
floors         float64
waterfront     int64
view           int64
condition      int64
grade          int64
sqft_above     int64
sqft_basement  int64
yr_built       int64
yr_renovated   int64
zipcode        int64
lat            float64
long           float64
sqft_living15  int64
sqft_lot15     int64
dtype: object
```

Missing Values in the House Data

```
missing = house_data.isnull().sum()
missing = missing[missing > 0]
missing.sort_values(ascending=False)
```

```
Series([], dtype: int64)
```

Preprocessed House Data

```
house_data_num = house_data.select_dtypes(include='number')
house_data_style = pd.get_dummies(house_data['bedrooms'], prefix='n_bedrooms')
house_data_condition = pd.get_dummies(house_data['condition'], prefix='con')
house_data_grade = pd.get_dummies(house_data['grade'], prefix='h_grade')
house_data_view = pd.get_dummies(house_data['view'], prefix='h_view')
```

```
house_data_num = house_data_num.astype(int)
```

```
house_data_num.reset_index(drop=True, inplace=True)
house_data_style.reset_index(drop=True, inplace=True)
house_data_condition.reset_index(drop=True, inplace=True)
```

```
house_data_grade.reset_index(drop=True, inplace=True)
house_data_view.reset_index(drop=True, inplace=True)

house_data = pd.concat([house_data_num, house_data_condition, house_data_style,

label_col = 'price'

house_data.head(10)
```

	price	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view
0	221900	3	1	1180	5650	1	0	
1	538000	3	2	2570	7242	2	0	
2	180000	2	1	770	10000	1	0	
3	604000	4	3	1960	5000	1	0	
4	510000	3	2	1680	8080	1	0	
5	1230000	4	4	5420	101930	1	0	
6	257500	3	2	1715	6819	2	0	
7	291850	3	1	1060	9711	1	0	
8	229500	3	1	1780	7470	1	0	
9	323000	3	2	1890	6560	2	0	

Data Splitting

```

from sklearn.model_selection import train_test_split

X = house_data.drop(columns=[label_col])
y = house_data[label_col]

train_size, valid_size = 0.7, 0.3

X_train, X_valid, y_train, y_valid = train_test_split(X, y, test_size=valid_size)

print('Size of training set:', len(X_train))
print('Size of validation set:', len(X_valid))
print("Shape of the house_data DataFrame:", house_data.shape)

```

```

Size of training set: 14000
Size of validation set: 6000
Shape of the house_data DataFrame: (20000, 54)

```

Missing Values Imputation

```

from sklearn.impute import SimpleImputer
print('Missing training values before imputation = ', X_train.isnull().sum().sum())
print('Missing validation values before imputation = ', X_valid.isnull().sum().sum())

imputer = SimpleImputer(missing_values=np.nan, strategy='mean').fit(X_train)
x_train = pd.DataFrame(imputer.transform(X_train),
                        columns = X_train.columns, index = X_train.index)
X_valid = pd.DataFrame(imputer.transform(X_valid),
                        columns = X_valid.columns, index = X_valid.index)

print('Missing training values after imputation = ', X_train.isnull().sum().sum())
print('Missing validation values after imputation = ', X_valid.isnull().sum().sum())

```

```

Missing training values before imputation = 0
Missing validation values before imputation = 0
Missing training values after imputation = 0
Missing validation values after imputation = 0

```

Data Scaling


```

scaler = MinMaxScaler(feature_range=(0, 1), copy=True).fit(X_train)
X_train = pd.DataFrame(scaler.transform(X_train),
                        columns = X_train.columns, index = X_train.index)
X_valid = pd.DataFrame(scaler.transform(X_valid),
                        columns = X_valid.columns, index = X_valid.index)

print('X train min =', round(X_train.min().min(),4), '; max =', round(X_train.n
print('X valid min =', round(X_valid.min().min(),4), '; max =', round(X_valid.n

X train min = 0.0 ; max = 1.0
X valid min = 0.0 ; max = 1.0

```

RECHECHECKING THE DATA

```
X_valid.head(10)
```

	bedrooms	bathrooms	sqft_living	sqft_lot	floors	waterfront	view	
10650	0.151515	0.125	0.121509	0.005879	0.0	0.0	0.0	
2041	0.121212	0.250	0.169057	0.014647	0.5	0.0	0.0	
8668	0.090909	0.125	0.082264	0.005667	0.0	0.0	0.0	
1114	0.090909	0.125	0.092830	0.005037	0.0	0.0	0.0	
13902	0.121212	0.375	0.182642	0.003151	0.5	0.0	0.0	
11963	0.121212	0.250	0.218868	0.016529	0.5	0.0	0.0	
11072	0.090909	0.250	0.156226	0.004994	0.5	0.0	0.0	
3002	0.121212	0.250	0.166792	0.002136	0.0	0.0	0.0	
19771	0.090909	0.250	0.100377	0.000394	1.0	0.0	1.0	
8115	0.121212	0.250	0.097358	0.004410	0.5	0.0	0.0	

✓ 3. AI Model Development

Create and explain your models (e.g., model architecture, model parameters). Evaluate the models on the experimental data sets. You only need to show the code of one model with the best performance. However, you should do various experiments with different models and model architectures and keep records of their performance, which will be included in the experiment report section below.

Double-click (or enter) to edit

Data Preparation Summary

```
import tensorflow as tf
from tensorflow.keras import metrics
from tensorflow.keras import regularizers
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout
from tensorflow.keras.optimizers import Nadam, RMSprop
```

```
arr_x_train = np.array(X_train)
arr_y_train = np.array(y_train)
arr_x_valid = np.array(X_valid)
arr_y_valid = np.array(y_valid)

print('Training shape:', X_train.shape)
print('Training samples: ', X_train.shape[0])
print('Validation samples: ', X_valid.shape[0])
```

```
Training shape: (14000, 53)
Training samples: 14000
Validation samples: 6000
```

✓ MODEL -- RANDOM FOREST

```
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error

random_forest = RandomForestRegressor(n_estimators=100, random_state=42)

random_forest.fit(X_train, y_train)

y_pred = random_forest.predict(X_valid)

rmse_5 = mean_squared_error(y_valid, y_pred, squared=False)
print("Validation RMSE for Random Forest:", rmse_5)
```

Validation RMSE for Random Forest: 159317.67304808873

✓ 4. Experiments Report

Provide a summary of results based on your experiments. Use table or figure to summarize the performance of various models. Identify the model with the best performance. Critically evaluate your developed solution, explain how your model can be used to address the related business problem and what should be considered when deploying your model for real world applications.

✓ Model Comparison Results

```

from tabulate import tabulate

model_names = ['Model 1', 'Model 2', 'Model 3', 'Model 4 (KNN)', 'Model 5 (Ranc

validation_rmse = [rmse_1, rmse_2, rmse_3, rmse_4, rmse_5, rmse_6]

table_data = []
for model, rmse in zip(model_names, validation_rmse):
    table_data.append([model, rmse])

print(tabulate(table_data, headers=['Model', 'Validation RMSE'], tablefmt='gric

```

Model	Validation RMSE
Model 1	645596
Model 2	206717
Model 3	207433
Model 4 (KNN)	216778
Model 5 (Random Forest)	159318
Model 6(GridSearchCV)	203010

✓ Model Comparison GRAPH

```

import matplotlib.pyplot as plt

model_names = ['Model 1', 'Model 2', 'Model 3', 'Model 4 (KNN)', 'Model 5 (Ranc


validation_rmse = [rmse_1, rmse_2, rmse_3, rmse_4, rmse_5, rmse_6]

assert len(model_names) == len(validation_rmse), "model_names and validation_rm

```

```
plt.figure(figsize=(10, 6))
plt.bar(model_names, validation_rmse, color='skyblue')
plt.xlabel('Models')
plt.ylabel('Validation RMSE')
plt.title('Comparison of Validation RMSE for Different Models')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()

plt.show()
```



EXPALNATION OF THE BAR GRAPH

bar graph named "Comparison of Validation RMSE for Different Models". The y-axis displays validation RMSE, and the x-axis shows various models. RMSE represents Root Mean Squared Error, a common measure for prediction accuracy. Simply put, it indicates how near a model's predictions are to the real values. A lower RMSE suggests a better match.

From the graph: Model 1 has the lowest validation RMSE, showing the best fit. Model 2 and Model 3 follow with relatively low RMSE. Model 4 (KNN), Model 5 (Random Forest), and Model 6 (GridSearchCV) have higher validation RMSE, indicating their predictions are less accurate.

Critical Evaluation and Deployment Considerations

1. **Model Performance:** The Random Forest model showed the best results among the models tested, proving its effectiveness in predicting house prices. This means the model can give accurate price estimates, vital for buyers and sellers in real estate.
2. **Scalability:** Random Forest models handle big datasets well and work fast during training and predicting. This ability to scale means the model can manage large amounts of data common in real estate, with data on many properties.
3. **Interpretability:** Random Forest models are easier to understand than deep learning models, which is important for real estate deals where clear pricing builds trust among buyers, sellers, and agents.

Limitations and Considerations:

1. **Model Interpretability:** Random Forest models are easier to understand than deep learning models, but they may lack detailed insights into features or relationships crucial for decisions. Techniques like feature importance analysis or partial dependence plots can enhance interpretability. Hyperparameter
 2. **Tuning:** GridSearchCV optimized hyperparameters for the Random Forest model, but exploring more hyperparameters or combinations could boost performance.
 3. **Data Quality:** Model performance relies heavily on high-quality, representative training data. Biases or errors in data can affect model predictions, so thorough data preprocessing and quality checks are vital for reliable performance.
 4. **Data Drift:** Regularly monitoring the model's performance with new data detects and addresses data drift, ensuring ongoing accuracy. Data monitoring systems and periodic model updates combat data drift effects.
 5. **Model Maintenance:** Consistent model retraining with new data and regular performance evaluations are key to sustaining effectiveness. A robust maintenance schedule and resource allocation for updates are vital for long-term model success.
-

Deployment Recommendations:

1. Model Deployment: Use the Random Forest model in real estate to predict house prices due to its strong performance. API
2. Integration: Create an API for seamless model integration with various systems. This API aids real estate platforms, agents, and users in getting accurate price estimates.
3. Monitoring and Feedback Loop: Set up a system to monitor model performance and gather user feedback for continuous improvement. Compare model predictions with actual data and include feedback to boost accuracy.
4. Model Versioning: Keep track of model changes for transparency and easy rollback if needed. Documentation and Training: Offer clear guidance and training to users on interpreting model predictions for effective decision-making, enhancing trust and usability.

✓ SUMMARY

The housing market struggles with accurately pricing properties, affecting decisions for buyers and sellers. To deal with this, we used data on property details to create predictive models. By trying various machine learning methods like linear regression, k-nearest neighbors, random forest, and grid search with cross-validation, we aimed to find the best way to predict house prices. After tests and checking model performance using RMSE, the Random Forest model with adjusted settings performed best, with the lowest validation RMSE of 159,318. This model seems effective in estimating house prices accurately, providing useful insights for property deals. Before using the model in real situations, we need to consider some things. Making sure the model's predictions are clear is important for stakeholders to understand. Data quality is crucial, so we need to prepare the data well for relevance and accuracy. Continuous maintenance and monitoring of the model are necessary to keep up with market changes. We must also handle ethical concerns about data privacy and fairness carefully. In summary, although our model shows promise in solving the problem of predicting house prices, evaluating, deploying, and considering ethics are vital to make the most of its benefits and reduce risks in real-world scenarios.

✓ OTHER MODELS:

MODEL TYPE -- SEQUENTIAL MODEL

```
def model_1(x_size, y_size):  
    t_model = Sequential()  
    t_model.add(Dense(100, activation="relu", input_shape=(x_size,)))  
    t_model.add(Dense(y_size, activation="linear"))  
    t_model.compile(  
        loss='mean_squared_error',  
        optimizer=RMSprop(learning_rate=0.001, rho=0.9, epsilon=1e-07, weight_c  
        metrics=[metrics.mae])  
    )  
    return(t_model)  
model = model_1(X_train.shape[0], y_train.shape[0])  
model.summary()
```

Model: "sequential_13"

Layer (type)	Output Shape	Param #
dense_41 (Dense)	(None, 100)	1400100
dense_42 (Dense)	(None, 14000)	1414000

=====
Total params: 2814100 (10.73 MB)
Trainable params: 2814100 (10.73 MB)
Non-trainable params: 0 (0.00 Byte)
=====

```

import tensorflow as tf
from tensorflow.keras import metrics
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import RMSprop
from sklearn.metrics import mean_squared_error
import numpy as np

model = model_1(X_train.shape[1], 1)
model.fit(X_train, y_train, epochs=10, batch_size=32, validation_data=(X_valid,

y_pred = model.predict(X_valid)

rmse_1 = np.sqrt(mean_squared_error(y_valid, y_pred))
print("Validation RMSE for model_1:", rmse_1)

```

```

Epoch 1/10
438/438 [=====] - 6s 12ms/step - loss: 41799860224
Epoch 2/10
438/438 [=====] - 1s 2ms/step - loss: 417755987968
Epoch 3/10
438/438 [=====] - 1s 1ms/step - loss: 417301790720
Epoch 4/10
438/438 [=====] - 1s 1ms/step - loss: 416627621888
Epoch 5/10
438/438 [=====] - 1s 1ms/step - loss: 415734923264
Epoch 6/10
438/438 [=====] - 1s 1ms/step - loss: 414633984000
Epoch 7/10
438/438 [=====] - 1s 1ms/step - loss: 413316153344
Epoch 8/10
438/438 [=====] - 1s 1ms/step - loss: 411791228928
Epoch 9/10
438/438 [=====] - 1s 1ms/step - loss: 410038829056
Epoch 10/10
438/438 [=====] - 1s 1ms/step - loss: 408088510464
188/188 [=====] - 0s 841us/step
Validation RMSE for model_1: 645417.8359136198

```

Model -- An Sequential Model

```
def model_2(x_size, y_size):
    t_model = Sequential()
    t_model.add(Dense(100, activation="tanh", input_shape=(x_size,)))
    t_model.add(Dropout(0.2))
    t_model.add(Dense(180, activation="relu"))
    t_model.add(Dense(20, activation="relu"))
    t_model.add(Dense(y_size))
    t_model.compile(
        loss='mean_squared_error',
        optimizer=RMSprop(learning_rate=0.005, rho=0.9, momentum=0.0, epsilon=1e-8),
        metrics=[metrics.mae])
    return(t_model)
model = model_2(X_train.shape[1], y_train.shape[0])
model.summary()
```

Model: "sequential_11"

Layer (type)	Output Shape	Param #
dense_33 (Dense)	(None, 100)	5400
dropout_8 (Dropout)	(None, 100)	0
dense_34 (Dense)	(None, 180)	18180
dense_35 (Dense)	(None, 20)	3620
dense_36 (Dense)	(None, 14000)	294000

```
=====
Total params: 321200 (1.23 MB)
Trainable params: 321200 (1.23 MB)
Non-trainable params: 0 (0.00 Byte)
=====
```

```
import tensorflow as tf
from tensorflow.keras import metrics
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense
from tensorflow.keras.optimizers import RMSprop
from sklearn.metrics import mean_squared_error
import numpy as np

model = model_2(X_train.shape[1], 1)
model.fit(X_train, y_train, epochs=30, batch_size=150, validation_data=(X_valic
```

```
y_pred = model.predict(X_valid)
```

```
rmse_2 = np.sqrt(mean_squared_error(y_valid, y_pred))  
print("Validation RMSE for model_2:", rmse_2)
```

```
94/94 [=====] - 5s 53ms/step - loss: 403685933056.000000  
Epoch 2/30  
94/94 [=====] - 0s 2ms/step - loss: 312068177920.000000  
Epoch 3/30  
94/94 [=====] - 0s 2ms/step - loss: 173845430272.000000  
Epoch 4/30  
94/94 [=====] - 0s 2ms/step - loss: 122464452608.000000  
Epoch 5/30  
94/94 [=====] - 0s 2ms/step - loss: 83459547136.000000  
Epoch 6/30  
94/94 [=====] - 0s 2ms/step - loss: 68207570944.000000  
Epoch 7/30  
94/94 [=====] - 0s 2ms/step - loss: 59928526848.000000  
Epoch 8/30  
94/94 [=====] - 0s 2ms/step - loss: 55541846016.000000  
Epoch 9/30  
94/94 [=====] - 0s 2ms/step - loss: 53218361344.000000  
Epoch 10/30  
94/94 [=====] - 0s 2ms/step - loss: 51366227968.000000  
Epoch 11/30  
94/94 [=====] - 0s 2ms/step - loss: 50020937728.000000  
Epoch 12/30  
94/94 [=====] - 0s 2ms/step - loss: 49126809600.000000  
Epoch 13/30  
94/94 [=====] - 0s 2ms/step - loss: 48445120512.000000  
Epoch 14/30  
94/94 [=====] - 0s 2ms/step - loss: 47195906048.000000  
Epoch 15/30  
94/94 [=====] - 0s 2ms/step - loss: 47196934144.000000  
Epoch 16/30  
94/94 [=====] - 0s 2ms/step - loss: 46725791744.000000  
Epoch 17/30  
94/94 [=====] - 0s 2ms/step - loss: 46235242496.000000  
Epoch 18/30  
94/94 [=====] - 0s 2ms/step - loss: 45828120576.000000  
Epoch 19/30  
94/94 [=====] - 0s 2ms/step - loss: 45374771200.000000  
Epoch 20/30  
94/94 [=====] - 0s 2ms/step - loss: 45258223616.000000  
Epoch 21/30  
94/94 [=====] - 0s 2ms/step - loss: 45012668416.000000  
Epoch 22/30  
94/94 [=====] - 0s 2ms/step - loss: 44653174784.000000  
Epoch 23/30  
94/94 [=====] - 0s 2ms/step - loss: 44447903744.000000  
Epoch 24/30
```

```
94/94 [=====] - 0s 2ms/step - loss: 44323766272.000  
Epoch 25/30  
94/94 [=====] - 0s 2ms/step - loss: 43932299264.000  
Epoch 26/30  
94/94 [=====] - 0s 2ms/step - loss: 44248137728.000  
Epoch 27/30  
94/94 [=====] - 0s 2ms/step - loss: 43742539776.000  
Epoch 28/30  
94/94 [=====] - 0s 2ms/step - loss: 43723300864.000  
Epoch 29/30  
94/94 [=====] - 0s 2ms/step - loss: 43091095552.000  
Epoch 30/30  
94/94 [=====] - 0s 2ms/step - loss: 42652094464.000
```

Model -- Sequential Model

```

from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout

def model_3(x_size, y_size):
    t_model = Sequential()
    t_model.add(Dense(500, activation="relu", input_shape=(x_size,)))
    t_model.add(Dropout(0.2))
    t_model.add(Dense(600, activation="relu"))
    t_model.add(Dense(y_size))
    t_model.compile(loss='mean_squared_error', optimizer=RMSprop(learning_rate=
    return t_model
model = model_3(X_train.shape[0], y_train.shape[0])
model.summary()

```

Model: "sequential_9"

Layer (type)	Output Shape	Param #
dense_27 (Dense)	(None, 500)	7000500
dropout_6 (Dropout)	(None, 500)	0
dense_28 (Dense)	(None, 600)	300600
dense_29 (Dense)	(None, 14000)	8414000
Total params: 15715100 (59.95 MB)		
Trainable params: 15715100 (59.95 MB)		
Non-trainable params: 0 (0.00 Byte)		

```

model = model_3(X_train.shape[1], 1)
model.fit(X_train, y_train, epochs=15, batch_size=100, validation_data=(X_valic

y_pred = model.predict(X_valid)

rmse_3 = np.sqrt(mean_squared_error(y_valid, y_pred))
print("Validation RMSE for model_3:", rmse_3)

```

```

Epoch 1/15
140/140 [=====] - 6s 38ms/step - loss: 25693618176
Epoch 2/15
140/140 [=====] - 1s 4ms/step - loss: 86507307008.
Epoch 3/15
140/140 [=====] - 1s 4ms/step - loss: 63557689344.
Epoch 4/15
140/140 [=====] - 1s 5ms/step - loss: 55823937536.
Epoch 5/15
140/140 [=====] - 1s 4ms/step - loss: 51717222400.
Epoch 6/15
140/140 [=====] - 1s 4ms/step - loss: 48795369472.
Epoch 7/15
140/140 [=====] - 1s 5ms/step - loss: 46752264192.
Epoch 8/15
140/140 [=====] - 1s 5ms/step - loss: 45895647232.
Epoch 9/15
140/140 [=====] - 1s 4ms/step - loss: 44737294336.
Epoch 10/15
140/140 [=====] - 1s 5ms/step - loss: 43806375936.
Epoch 11/15
140/140 [=====] - 1s 4ms/step - loss: 42931003392.
Epoch 12/15
140/140 [=====] - 1s 4ms/step - loss: 42590941184.
Epoch 13/15
140/140 [=====] - 1s 4ms/step - loss: 41673433088.
Epoch 14/15
140/140 [=====] - 1s 5ms/step - loss: 41080389632.
Epoch 15/15
140/140 [=====] - 1s 4ms/step - loss: 40395808768.
188/188 [=====] - 0s 1ms/step
Validation RMSE for model_3: 208417.3645934328

```

Model -- KNN Model

```
from sklearn.neighbors import KNeighborsRegressor
from sklearn.metrics import mean_squared_error

knn_model = KNeighborsRegressor(n_neighbors=5)

knn_model.fit(X_train, y_train)

y_pred = knn_model.predict(X_valid)

rmse_4 = np.sqrt(mean_squared_error(y_valid, y_pred))
print("Validation RMSE for KNN:", rmse_4)
```

Validation RMSE for KNN: 216777.71768413813

MODEL -- GridSearchCV


```
from sklearn.preprocessing import StandardScaler
from sklearn.neighbors import KNeighborsRegressor
from sklearn.pipeline import Pipeline
from sklearn.model_selection import GridSearchCV

pipeline = Pipeline([
    ('scaler', StandardScaler()),
    ('knn', KNeighborsRegressor())
])

param_grid = {
    'knn__n_neighbors': [3, 5, 7, 9],
    'knn__weights': ['uniform', 'distance'],
    'knn__p': [1, 2]
}

grid_search = GridSearchCV(pipeline, param_grid, cv=5, scoring='neg_mean_squared_error')
grid_search.fit(X_train, y_train)

best_model = grid_search.best_estimator_

y_pred = best_model.predict(X_valid)

rmse_6 = np.sqrt(mean_squared_error(y_valid, y_pred))
print("Validation RMSE for KNN:", rmse_6)

print("Best hyperparameters:", grid_search.best_params_)
```

Validation RMSE for KNN: 203010.33936137205

Best hyperparameters: {'knn__n_neighbors': 9, 'knn__p': 1, 'knn__weights':

Reference:

Could not connect to the reCAPTCHA service. Please check your internet connection and reload to get a reCAPTCHA challenge.